

# Numerical methods project

---

## HUBER-ENERGY VARIATIONAL AUTOENCODER

Grégoire MOURRE

David PREMACHANDRA

---



Master 2 Applied and Theoretical Mathematics

2023-2024

# Contents

1	First test with the MNIST dataset (convolutional encoder and decoder)	2
2	Second test with the MNIST dataset (dense encoder and decoder)	4
3	Third test with the MNIST dataset (dense encoder and decoder with more neurons)	6
4	Fourth test with the MNIST dataset (same parameters than in the previous question, but <code>lambda_factor = 0</code> )	8
5	Test with the Fashion-MNIST dataset	9
6	Test with the CIFAR-10 dataset	11
7	Test with the Boston Housing dataset (1D data)	13

## Abstract

In this file, we test several architectures for some well known datasets. Each time, we first present the parameters and the architecture of the encoder and decoder. Then, we plot the results.

# 1 First test with the MNIST dataset (convolutional encoder and decoder)

## Encoder Architecture

```
1 mnist_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['input_dim']),
3     Conv2D(32, (3, 3), activation='relu'),
4     Conv2D(64, (3, 3), activation='relu'),
5     Flatten(),
6     Dense(128, activation='relu'),
7     Dense(mnist_parameters['latent_dim'], activation=None)
8 ])
```


## Decoder Architecture

```
1 mnist_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['latent_dim']),
3     Dense(256, activation='relu'),
4     Dense(1024, activation='relu'),
5     Dense(8*8*8, activation='relu'),
6     Reshape((8,8,8)),
7     Conv2DTranspose(16, (8, 8), activation='relu'),
8     Conv2DTranspose(4, (8, 8), activation='relu'),
9     Conv2DTranspose(1, (7, 7), activation='sigmoid')
10 ])
```

## Model Parameters

```
1 mnist_parameters = {'input_dim': (28,28,1),
2                     '1D data': False,
3                     'latent_dim': 12}
4
5 mnist_training_parameters = {'dataset': train_images_mnist / 255.0,
6                             'learning_rate': 0.0005,
7                             'epochs': 2000,
8                             'batch_size': 200}
```


## Results




(a) Loss during the training



(b) Images and their reconstructed version



(c) Generated images



(d) Smooth transition between two images

Figure 1: Test of the architecture and the parameters presented above for the MNIST dataset

## 2 Second test with the MNIST dataset (dense encoder and decoder)

```
1 number_of_neurons = 128
```

### Encoder Architecture

```
1 mnist_2_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['input_dim']),
3     Flatten(),
4     Dense(number_of_neurons, activation='relu'),
5     Dense(number_of_neurons, activation='relu'),
6     Dense(number_of_neurons, activation='relu'),
7     Dense(mnist_parameters['latent_dim'], activation=None)
8 ])
```


### Decoder Architecture

```
1 mnist_2_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['latent_dim']),
3     Dense(number_of_neurons, activation='relu'),
4     Dense(number_of_neurons, activation='relu'),
5     Dense(28*28*1, activation='sigmoid'),
6     Reshape((28,28,1))
7 ])
```

### Model Parameters

```
1 mnist_parameters = {'input_dim': (28,28,1),
2                     '1D data': False,
3                     'latent_dim': 12}
4
5 mnist_training_parameters = {'dataset': train_images_mnist / 255.0,
6                              'learning_rate': 0.0005,
7                              'epochs': 2000,
8                              'batch_size': 200}
```


### Results




(a) Loss during the training



(b) Images and their reconstructed version



(c) Generated images



(d) Smooth transition between two images

Figure 2: Test of the architecture and the parameters presented above for the MNIST dataset

### 3 Third test with the MNIST dataset (dense encoder and decoder with more neurons)

```
1 number_of_neurons = 256
```

#### Encoder Architecture

```
1 mnist_2_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['input_dim']),
3     Flatten(),
4     Dense(number_of_neurons, activation='relu'),
5     Dense(number_of_neurons, activation='relu'),
6     Dense(number_of_neurons, activation='relu'),
7     Dense(mnist_parameters['latent_dim'], activation=None)
8 ])
```


#### Decoder Architecture

```
1 mnist_2_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = mnist_parameters['latent_dim']),
3     Dense(number_of_neurons, activation='relu'),
4     Dense(number_of_neurons, activation='relu'),
5     Dense(28*28*1, activation='sigmoid'),
6     Reshape((28,28,1))
7     ])
```

#### Model Parameters

```
1 mnist_parameters = {'input_dim': (28,28,1),
2                     '1D data': False,
3                     'latent_dim': 12}
4
5 mnist_training_parameters = {'dataset': train_images_mnist / 255.0,
6                              'learning_rate': 0.0005,
7                              'epochs': 2000,
8                              'batch_size': 200}
```


#### Results




(a) Loss during the training



(b) Images and their reconstructed version



(c) Generated images



(d) Smooth transition between two images


Figure 3: Test of the architecture and the parameters presented above for the MNIST dataset



#### 4 Fourth test with the MNIST dataset (same parameters than in the previous question, but $\lambda_{\text{factor}} = 0$ )

Setting  $\lambda_{\text{factor}}$  to 0 means that we ignore the latent loss.


#### Results




(a) Loss during the training



(b) Images and their reconstructed version



(c) Generated images



(d) Smooth transition between two images

Figure 4: Test of the architecture and the parameters presented above for the MNIST dataset

## 5 Test with the Fashion-MNIST dataset

### Encoder Architecture

```
1 fashion_mnist_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = fashion_mnist_parameters['input_dim']),
3     Conv2D(32, (3, 3), activation='relu'),
4     Conv2D(64, (3, 3), activation='relu'),
5     Flatten(),
6     Dense(128, activation='relu'),
7     Dense(fashion_mnist_parameters['latent_dim'], activation=None)
8 ])
```


### Decoder Architecture

```
1 fashion_mnist_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = fashion_mnist_parameters['latent_dim']),
3     Dense(256, activation='relu'),
4     Dense(1024, activation='relu'),
5     Dense(8*8*8, activation='relu'),
6     Reshape((8,8,8)),
7     Conv2DTranspose(16, (8, 8), activation='relu'),
8     Conv2DTranspose(4, (8, 8), activation='relu'),
9     Conv2DTranspose(1, (7, 7), activation='sigmoid')
10 ])
```


### Model Parameters

```
1 fashion_mnist_parameters = {
2     'input_dim': (28, 28, 1),
3     '1D_data': False,
4     'latent_dim': 12
5 }
6
7 fashion_mnist_training_parameters = {
8     'dataset': train_images_fashion_mnist / 255.0,
9     'learning_rate': 0.0005,
10    'epochs': 1000,
11    'batch_size': 200
12 }
```


### Results




(a) Loss during the training



(c) Generated images



(b) Images and their reconstructed version



(d) Smooth transition between two images

Figure 5: Test of the architecture and the parameters presented above for the Fashion-MNIST dataset

## 6 Test with the CIFAR-10 dataset

### Encoder Architecture

```
1 cifar10_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = cifar10_parameters['input_dim']),
3     Conv2D(3, (2, 2), activation='relu'),
4     Conv2D(32, (2, 2), strides=(2,2), activation='relu'),
5     Conv2D(32, (3, 3), activation='relu'),
6     Conv2D(32, (3, 3), activation='relu'),
7     Flatten(),
8     Dense(128, activation='relu'),
9     Dense(cifar10_parameters['latent_dim'], activation=None)
10 ])
```


### Decoder Architecture

```
1 cifar10_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = cifar10_parameters['latent_dim']),
3     Dense(128, activation='relu'),
4     Dense(16*16*32, activation='relu'),
5     Reshape((16,16,32)),
6     Conv2DTranspose(32, (2, 2), padding = 'same', activation='
7     relu'),
8     Conv2DTranspose(32, (2, 2), padding = 'same', activation='
9     relu'),
10    Conv2DTranspose(32, (3, 3), strides = (2,2), activation='
    relu'),
    Conv2D(3, (2, 2), activation='sigmoid')
])
```


### Model Parameters

```
1 cifar10_parameters = {'input_dim': train_images_cifar10[0].shape,
2                        '1D data': False,
3                        'latent_dim': 70}
4
5 cifar10_training_parameters = {'dataset': train_images_cifar10 / 255.0,
6                                'learning_rate': 0.0005,
7                                'epochs': 1000,
8                                'batch_size': 200}
```

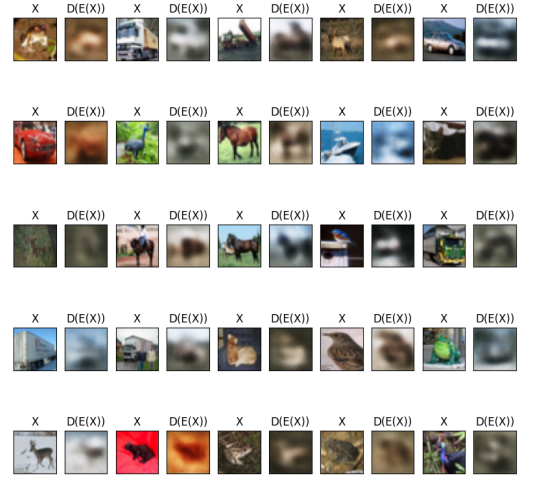
### Results




(a) Loss during the training



(c) Generated images



(b) Images and their reconstructed version



(d) Smooth transition between two images

Figure 6: Test of the architecture and the parameters presented above for the CIFAR-10 dataset

## 7 Test with the Boston Housing dataset (1D data)

### Encoder Architecture

```
1 boston_housing_encoder = tf.keras.Sequential([
2     InputLayer(input_shape = boston_housing_parameters['
   input_dim']),
3     Dense(256, activation='relu'),
4     Dense(256, activation='relu'),
5     Dense(256, activation='relu'),
6     Dense(boston_housing_parameters['latent_dim'], activation=
   None)
7 ])
```


### Decoder Architecture

```
1 boston_housing_decoder = tf.keras.Sequential([
2     InputLayer(input_shape = boston_housing_parameters['
   latent_dim']),
3     Dense(256, activation='relu'),
4     Dense(256, activation='relu'),
5     Dense(256, activation='relu'),
6     Dense(boston_housing_parameters['input_dim'], activation=
   None)
7 ])
```


### Model Parameters

```
1 boston_housing_parameters = {'input_dim': (train_data_boston_housing[0].
   shape)[0],
2                               '1D data': True,
3                               'latent_dim': 3}
4
5 boston_housing_training_parameters = {'dataset': train_data_boston_housing,
6                                       'learning_rate': 0.00005,
7                                       'epochs': 2000,
8                                       'batch_size': 300}
```

### Results




(a) Loss during the training



(c) Generated data



(b) Initial data and their reconstructed version



(d) Smooth transition between two data

Figure 7: Test of the architecture and the parameters presented above for 1D data of the Boston Housing data set